

02610

# Optimization and Data Fitting

## Week 11: Conjugate Gradient Methods & Large-Scale Unconstrained Optimization

Yiqiu Dong

DTU Compute  
Technical University of Denmark

16 November 2020

# Unconstrained quadratic problems

$$\min_{\mathbf{x} \in \mathbb{R}^n} f(\mathbf{x}) = \frac{1}{2} \mathbf{x}^T A \mathbf{x} - \mathbf{b}^T \mathbf{x}$$

where  $A$  is an  $n \times n$  symmetric positive definite matrix.

- It is equivalent to solve the linear system of equations:  $A\mathbf{x} = \mathbf{b}$ .
- The **residual**  $\mathbf{r} = \mathbf{b} - A\mathbf{x}$  is the negative gradient:  $\mathbf{r} = -\nabla f(\mathbf{x})$ .
- Steepest descent method: 1 or  $\infty$  iterations.
- Newton's method: 1 iteration.
- Coordinate search method:  $n$  or  $\infty$  iterations.
- Conjugate gradient method:  $n$  iterations.

# Unconstrained quadratic problems

$$\min_{\mathbf{x} \in \mathbb{R}^n} f(\mathbf{x}) = \frac{1}{2} \mathbf{x}^T A \mathbf{x} - \mathbf{b}^T \mathbf{x}$$

where  $A$  is an  $n \times n$  symmetric positive definite matrix.

- It is equivalent to solve the linear system of equations:  $A\mathbf{x} = \mathbf{b}$ .
- The **residual**  $\mathbf{r} = \mathbf{b} - A\mathbf{x}$  is the negative gradient:  $\mathbf{r} = -\nabla f(\mathbf{x})$ .
- **Steepest descent method:** 1 or  $\infty$  iterations.
- **Newton's method:** 1 iteration.
- **Coordinate search method:**  $n$  or  $\infty$  iterations.
- **Conjugate gradient method:**  $n$  iterations.

# Conjugate gradient (CG) method

- It was proposed by Hestenes and Stiefel in the 1950s.
- It is the most widely used iterative method for solving  $A\mathbf{x} = \mathbf{b}$  with  $A \succ 0$
- It was extended to solve nonlinear unconstrained minimization problems in 1960s.



Main **advantages** of CG:

- It takes at most  $n$  iterations to the solution (theoretically).
- It does not alter  $A$ .
- At each iteration, it only need one computation of the matrix-vector product ( $O(n^2)$ ) and a few vector product and sum ( $O(n)$ ).
- For storage, it only need store a few vectors.
- CG is only used for solving large-scale problems.
- CG is proved with linear convergence rate, but generally much faster than the steepest descent method.

# Conjugate directions

## Idea:

- If  $A$  is diagonal, then the coordinate search method can find the minimizer of  $f(\mathbf{x})$  in  $n$  iterations.





- If  $A$  is **NOT** diagonal, we can diagonalize  $A$ , that is, accordingly transform the coordinate directions.

# Conjugate directions

## Idea:

- If  $A$  is diagonal, then the coordinate search method can find the minimizer of  $f(\mathbf{x})$  in  $n$  iterations.



- If  $A$  is **NOT** diagonal, we can diagonalize  $A$ , that is, accordingly transform the coordinate directions.

## Conjugate directions

Suppose that a  $n \times n$  matrix  $S = [\mathbf{p}_0, \mathbf{p}_1, \dots, \mathbf{p}_{n-1}]$  diagonalizes  $A$ , i.e.,  $S^T A S$  is diagonal. Then, we have

$$\mathbf{p}_i^T A \mathbf{p}_j = 0, \quad \text{for all } i \neq j,$$

and we call  $\{\mathbf{p}_0, \dots, \mathbf{p}_{n-1}\}$  to be **conjugate** with respect to spd.  $A$ .

- $\{\mathbf{p}_0, \dots, \mathbf{p}_{n-1}\}$  are conjugate, if and only if they are orthogonal for the inner product  $\langle \mathbf{u}, \mathbf{v} \rangle_A = \mathbf{u}^T A \mathbf{v}$ .
- If  $\mathbf{p}_i \neq \mathbf{0}$  for all  $i$ , they are also linearly independent.

If  $\{\mathbf{p}_0, \dots, \mathbf{p}_{n-1}\}$  are **conjugate** (*conjugate directions*) and  $\mathbf{x}_{k+1} = \mathbf{x}_k + \alpha_k \mathbf{p}_k$ , then the exact line search has a closed-form and gives

$$\alpha_k = \frac{\mathbf{r}_k^T \mathbf{p}_k}{\mathbf{p}_k^T A \mathbf{p}_k}.$$

## Conjugate directions

Suppose that a  $n \times n$  matrix  $S = [\mathbf{p}_0, \mathbf{p}_1, \dots, \mathbf{p}_{n-1}]$  diagonalizes  $A$ , i.e.,  $S^T A S$  is diagonal. Then, we have

$$\mathbf{p}_i^T A \mathbf{p}_j = 0, \quad \text{for all } i \neq j,$$

and we call  $\{\mathbf{p}_0, \dots, \mathbf{p}_{n-1}\}$  to be **conjugate** with respect to spd.  $A$ .

- $\{\mathbf{p}_0, \dots, \mathbf{p}_{n-1}\}$  are conjugate, if and only if they are orthogonal for the inner product  $\langle \mathbf{u}, \mathbf{v} \rangle_A = \mathbf{u}^T A \mathbf{v}$ .
- If  $\mathbf{p}_i \neq \mathbf{0}$  for all  $i$ , they are also linearly independent.

If  $\{\mathbf{p}_0, \dots, \mathbf{p}_{n-1}\}$  are **conjugate** (*conjugate directions*) and  $\mathbf{x}_{k+1} = \mathbf{x}_k + \alpha_k \mathbf{p}_k$ , then the exact line search has a closed-form and gives

$$\alpha_k = \frac{\mathbf{r}_k^T \mathbf{p}_k}{\mathbf{p}_k^T A \mathbf{p}_k}.$$



## Conjugate directions

Suppose that a  $n \times n$  matrix  $S = [\mathbf{p}_0, \mathbf{p}_1, \dots, \mathbf{p}_{n-1}]$  diagonalizes  $A$ , i.e.,  $S^T A S$  is diagonal. Then, we have

$$\mathbf{p}_i^T A \mathbf{p}_j = 0, \quad \text{for all } i \neq j,$$

and we call  $\{\mathbf{p}_0, \dots, \mathbf{p}_{n-1}\}$  to be **conjugate** with respect to spd.  $A$ .

- $\{\mathbf{p}_0, \dots, \mathbf{p}_{n-1}\}$  are conjugate, if and only if they are orthogonal for the inner product  $\langle \mathbf{u}, \mathbf{v} \rangle_A = \mathbf{u}^T A \mathbf{v}$ .
- If  $\mathbf{p}_i \neq \mathbf{0}$  for all  $i$ , they are also linearly independent.

If  $\{\mathbf{p}_0, \dots, \mathbf{p}_{n-1}\}$  are **conjugate** (*conjugate directions*) and  $\mathbf{x}_{k+1} = \mathbf{x}_k + \alpha_k \mathbf{p}_k$ , then the exact line search has a closed-form and gives

$$\alpha_k = \frac{\mathbf{r}_k^T \mathbf{p}_k}{\mathbf{p}_k^T A \mathbf{p}_k}.$$

# Conjugate direction methods

## Algorithm

Given  $\mathbf{x}_0$  and a set of conjugate directions  $\{\mathbf{p}_0, \dots, \mathbf{p}_{n-1}\}$ .

**loop**

    Compute  $\alpha_k = \frac{\mathbf{r}_k^T \mathbf{p}_k}{\mathbf{p}_k^T A \mathbf{p}_k};$

    Update  $\mathbf{x}_{k+1} = \mathbf{x}_k + \alpha_k \mathbf{p}_k;$

**end loop**

## Theorem

For any  $\mathbf{x}_0 \in \mathbb{R}^n$  the sequence  $\{\mathbf{x}_n\}$  generated by the above conjugate direction method converges to the solution  $\mathbf{x}^*$  of the linear system  $A\mathbf{x} = \mathbf{b}$  in at most  $n$  iterations.

# Conjugate direction methods

## Expanding subspace minimization

Let  $\mathbf{x}_0 \in \mathbb{R}^n$  be any starting point and the sequence  $\{\mathbf{x}_n\}$  be generated by the conjugate direction method shown in the previous page. Then,

- $\mathbf{r}_k^T \mathbf{p}_i = 0$ , for  $i = 0, 1, \dots, k-1$ ;
- $\mathbf{x}_k$  is the minimizer of  $f(\mathbf{x})$  over the set  $\{\mathbf{x} | \mathbf{x} = \mathbf{x}_0 + \text{span}\{\mathbf{p}_0, \dots, \mathbf{p}_{k-1}\}\}$ .

- The current residual  $\mathbf{r}_k$  is orthogonal to all previous search directions.
- The conjugate direction method minimizes  $f(\mathbf{x})$  along one conjugate direction at one iteration.

# Conjugate gradient directions

- $\mathbf{p}_k$  is generated by using only the previous vector  $\mathbf{p}_{k-1}$ .
- $\mathbf{p}_k$  is automatically conjugate to  $\{\mathbf{p}_0, \dots, \mathbf{p}_{k-1}\}$ .

**Recursion for  $\mathbf{p}_k$ :** We start with  $\mathbf{p}_0 = \mathbf{r}_0$  and choose  $\mathbf{p}_k$  to be a linear combination of the residual  $\mathbf{r}_k$  and the previous direction  $\mathbf{p}_{k-1}$ :

$$\mathbf{p}_k = \mathbf{r}_k + \beta_k \mathbf{p}_{k-1}.$$

Since  $\mathbf{p}_k$  is conjugate to  $\mathbf{p}_{k-1}$  w.r.t.  $A$ , then we have

$$\beta_k = -\frac{\mathbf{p}_{k-1}^T A \mathbf{r}_k}{\mathbf{p}_{k-1}^T A \mathbf{p}_{k-1}}.$$

# Conjugate gradient method (preliminary version)

## Algorithm

Given  $\mathbf{x}_0$ ;

Set  $\mathbf{r}_0 = \mathbf{b} - A\mathbf{x}_0$ ,  $\mathbf{p}_0 = \mathbf{r}_0$ ;

**loop**

    Compute  $\alpha_k = \frac{\mathbf{r}_k^T \mathbf{p}_k}{\mathbf{p}_k^T A \mathbf{p}_k}$ ;

    Update  $\mathbf{x}_{k+1} = \mathbf{x}_k + \alpha_k \mathbf{p}_k$ ;

    Compute  $\mathbf{r}_{k+1} = \mathbf{b} - A\mathbf{x}_{k+1}$ ;

    Compute  $\beta_{k+1} = -\frac{\mathbf{p}_k^T A \mathbf{r}_{k+1}}{\mathbf{p}_k^T A \mathbf{p}_k}$ ;

    Compute  $\mathbf{p}_{k+1} = \mathbf{r}_{k+1} + \beta_{k+1} \mathbf{p}_k$ ;

    Check for convergence;

**end loop**

Output  $\mathbf{x}_{k+1}$ .

# Properties of CG method

## Theorem

Suppose that the  $k$ th iterate of the CG method is not the solution  $\mathbf{x}^*$ . Then,

- ①  $\mathbf{r}_k^T \mathbf{r}_i = 0$ , for  $i = 0, 1, \dots, k-1$ ,
- ②  $\text{span}\{\mathbf{r}_0, \mathbf{r}_1, \dots, \mathbf{r}_k\} = \text{span}\{\mathbf{p}_0, \mathbf{p}_1, \dots, \mathbf{p}_k\} = \text{span}\{\mathbf{r}_0, A\mathbf{r}_0, \dots, A^k \mathbf{r}_0\}$ ,
- ③  $\mathbf{p}_k^T A \mathbf{p}_i = 0$ , for  $i = 0, 1, \dots, k-1$ .

Therefore, the sequence  $\{\mathbf{x}_k\}$  converges to  $\mathbf{x}^*$  in at most  $n$  steps.

- The proof of this theorem relies on the fact that  $\mathbf{p}_0 = \mathbf{r}_0$  (the steepest descent direction).
- The result (1) shows that the residuals/gradients at all iterates are orthogonal to each other.
- The result (4) shows that  $\{\mathbf{p}_0, \dots, \mathbf{p}_k\}$  are conjugate directions.
- The result (3) shows that the search directions and the residuals from CG method generate the Krylov subspaces.

# Krylov subspaces

**Definition:** For a linear system  $A\mathbf{x} = \mathbf{b}$ , a sequence of subspaces

$$\mathcal{K}_0 = \{\mathbf{0}\}, \quad \mathcal{K}_k = \text{span}\{\mathbf{b}, A\mathbf{b}, \dots, A^{k-1}\mathbf{b}\} \quad \text{for } k \geq 1.$$

## Properties:

- $\mathcal{K}_k = \text{span}\{\mathbf{r}_0, A\mathbf{r}_0, \dots, A^k\mathbf{r}_0\}$ .
- The Krylov subspaces are nested:  $\mathcal{K}_0 \subseteq \mathcal{K}_1 \subseteq \mathcal{K}_2 \subseteq \dots$
- The dimensions of the Krylov subspaces increase by at most one:  $\dim\mathcal{K}_{k+1} - \dim\mathcal{K}_k$  is zero or one.
- If  $\mathcal{K}_{k+1} = \mathcal{K}_k$ , then  $\mathcal{K}_i = \mathcal{K}_k$  for all  $i \geq k$ :

$$\begin{aligned} A^k\mathbf{b} &\in \text{span}\{\mathbf{b}, A\mathbf{b}, \dots, A^{k-1}\mathbf{b}\} \\ \implies A^i\mathbf{b} &\in \text{span}\{\mathbf{b}, A\mathbf{b}, \dots, A^{k-1}\mathbf{b}\} \quad \text{for } i > k. \end{aligned}$$

# Simplified CG method

- Using  $\mathbf{x}_{k+1} = \mathbf{x}_k + \alpha_k \mathbf{p}_k$  and  $\mathbf{r}_{k+1} = \mathbf{b} - A\mathbf{x}_{k+1}$ , we obtain

$$\mathbf{r}_{k+1} = \mathbf{r}_k - \alpha_k A\mathbf{p}_k. \quad (1)$$

- Using  $\mathbf{p}_k = \mathbf{r}_k + \beta_k \mathbf{p}_{k-1}$  and  $\mathbf{r}_k^T \mathbf{p}_{k-1} = 0$ , we obtain  $\mathbf{r}_k^T \mathbf{p}_k = \mathbf{r}_k^T \mathbf{r}_k$ , then

$$\alpha_k = \frac{\mathbf{r}_k^T \mathbf{p}_k}{\mathbf{p}_k^T A \mathbf{p}_k} = \frac{\|\mathbf{r}_k\|_2^2}{\mathbf{p}_k^T A \mathbf{p}_k}. \quad (2)$$

- Using (1), (2) and  $\mathbf{r}_{k+1}^T \mathbf{r}_k = 0$ , we obtain

$$\beta_{k+1} = -\frac{\mathbf{p}_k^T A \mathbf{r}_{k+1}}{\mathbf{p}_k^T A \mathbf{p}_k} = \frac{\|\mathbf{r}_{k+1}\|_2^2}{\|\mathbf{r}_k\|_2^2}. \quad (3)$$



# Conjugate gradient method

## Algorithm

Given  $\mathbf{x}_0$ ;

Set  $\mathbf{r}_0 = \mathbf{b} - A\mathbf{x}_0$ ,  $\mathbf{p}_0 = \mathbf{r}_0$ ;

**loop**

    Compute  $\alpha_k = \frac{\|\mathbf{r}_k\|_2^2}{\mathbf{p}_k^T A \mathbf{p}_k}$ ;

    Update  $\mathbf{x}_{k+1} = \mathbf{x}_k + \alpha_k \mathbf{p}_k$ ;

    Compute  $\mathbf{r}_{k+1} = \mathbf{r}_k - \alpha_k A \mathbf{p}_k$ ;

    Compute  $\beta_{k+1} = \frac{\|\mathbf{r}_{k+1}\|_2^2}{\|\mathbf{r}_k\|_2^2}$ ;

    Compute  $\mathbf{p}_{k+1} = \mathbf{r}_{k+1} + \beta_{k+1} \mathbf{p}_k$ ;

    Check for convergence;

**end loop**

Output  $\mathbf{x}_{k+1}$ .


Main computation per iteration is matrix-vector product  $A\mathbf{p}_{k+1}$ .

# Rate of convergence

- If  $A$  has only  $r$  distinct eigenvalues, then the CG method will terminate at the solution in at most  $r$  iterations.
- If  $A$  has eigenvalues  $\lambda_1 \leq \lambda_2 \leq \dots \leq \lambda_n$ , we have that

$$\|\mathbf{x}_{k+1} - \mathbf{x}^*\|_A^2 \leq \left( \frac{\lambda_{n-k} - \lambda_1}{\lambda_{n-k} + \lambda_1} \right)^2 \|\mathbf{x}_0 - \mathbf{x}^*\|_A^2.$$

**Example:** We apply the CG method to solve  $A\mathbf{x} = \mathbf{b}$ .



# Preconditioning

- **Idea:** Make change of variables  $\hat{\mathbf{x}} = C\mathbf{x}$  with  $C$  nonsingular, and apply CG to

$$C^{-T}AC^{-1}\hat{\mathbf{x}} = C^{-T}\mathbf{b}.$$

- The spectrum of the new matrix  $C^{-T}AC^{-1}$  should be clustered, then PCG converges fast.
- We need consider the trade-off between enhanced convergence and cost of extra computation.
- The matrix  $M = C^T C$  is called the **preconditioner**.
- Matlab implementation: `pcg`

## Example:

- diagonal  $C = \text{diag}(A_{11}, A_{22}, \dots, A_{nn})$
- incomplete or approximate Cholesky factorization of  $A$
- Good preconditioners are often application-dependent.

# Nonlinear conjugate gradient method

$$\min_{\mathbf{x} \in \mathbb{R}^n} f(\mathbf{x}), \quad f \text{ is convex and differentiable.}$$

## Nonlinear CG methods

- Extend linear CG method to nonquadratic functions.
- Limited global convergence theory.

**Modifications** needed to extend linear CG method

- Replace  $\mathbf{r}_k = \mathbf{b} - A\mathbf{x}_k$  with  $-\nabla f(\mathbf{x}_k)$ .
- Determine the step length  $\alpha$  by line search.

# Fletcher-Reeves method

## Algorithm

Given  $\mathbf{x}_0$ ;

Compute  $f_0 = f(\mathbf{x}_0)$  and  $\nabla f_0 = \nabla f(\mathbf{x}_0)$ ;

Set  $\mathbf{p}_0 = -\nabla f_0$ ;

**loop**

    Compute  $\alpha_k$  by line search method;

    Update  $\mathbf{x}_{k+1} = \mathbf{x}_k + \alpha_k \mathbf{p}_k$ ;

    Evaluate  $\nabla f_{k+1}$ ;

    Compute  $\beta_{k+1}^{FR} = \frac{\|\nabla f_{k+1}\|_2^2}{\|\nabla f_k\|_2^2}$ ;

    Compute  $\mathbf{p}_{k+1} = -\nabla f_{k+1} + \beta_{k+1}^{FR} \mathbf{p}_k$ ;

    Check for convergence;

**end loop**

Output  $\mathbf{x}_{k+1}$ .

# Some observations

## Interpretation

- First iteration is a steepest descent step.
- General update is a steepest descent step with momentum term

$$\mathbf{x}_{k+1} = \mathbf{x}_k - \alpha_k \nabla f_k + \frac{\alpha_k \beta_k}{\alpha_{k-1}} (\mathbf{x}_k - \mathbf{x}_{k-1}).$$

- It is common to restart the algorithm every  $n$  iterations by taking a steepest descent step to periodically refresh the algorithm.

## Line search

- With exact line search, it reduces to linear CG for quadratic  $f$ .
- Exact line search in computation of  $\alpha_k$  implies that  $\alpha_k$  is a local minimizer along  $\mathbf{p}_k$ , i.e.,  $\nabla f_{k+1}^T \mathbf{p}_k = 0$ . Therefore,  $\mathbf{p}_{k+1}$  is a descent direction at  $\mathbf{x}_{k+1}$ :

$$\nabla f_{k+1}^T \mathbf{p}_{k+1} = -\|\nabla f_{k+1}\|^2 + \beta_{k+1}^{FR} \nabla f_{k+1}^T \mathbf{p}_k = -\|\nabla f_{k+1}\|^2 < 0.$$

- For inexact line search, if  $\alpha_k$  satisfies the strong Wolfe conditions, then  $\mathbf{p}_{k+1}$  is descent.

## Variations

**Polak-Ribière method:** Compute  $\beta_{k+1}$  from

$$\beta_{k+1} = \frac{\nabla f_{k+1}^T (\nabla f_{k+1} - \nabla f_k)}{\|\nabla f_k\|_2^2}.$$

**Hestenes-Stiefel method:** Compute  $\beta_{k+1}$  from

$$\beta_{k+1} = \frac{\nabla f_{k+1}^T (\nabla f_{k+1} - \nabla f_k)}{(\nabla f_{k+1} - \nabla f_k)^T \mathbf{p}_k}.$$

- All these formulas are equivalent for quadratic  $f$  and exact line search.
- With restarts and the strong Wolfe conditions, all three methods have global convergence.
- Without restarts, FR has global convergence with the strong Wolfe conditions, but PR not.
- In practice, PR is more robust and efficient than FR.

# Large-scale unconstrained optimization

- Large-scale problems (today):  $10^3 \sim 10^6$  variables.
- When solving large-scale problems, we have to take the storage and computational costs of the optimization algorithm into account.
- In large problems, the following can have a prohibitive cost:
  - ▶ factorizing the Hessian (solving for the Newton step)
  - ▶ computing the Hessian or multiplying it
  - ▶ storing a dense approximate Hessian like in quasi-Newton methods
- Linear/nonlinear conjugate gradient methods can be applied directly to large-scale problems without modification, but not fast.



# Inexact Newton methods

**Ideas:** Use some inexpensive iterative algorithm to *very approximately* solve either

$$\nabla^2 f_k \mathbf{p}_k = -\nabla f_k \quad (\text{line search})$$

or

$$\begin{aligned} \min_{\mathbf{p} \in \mathbb{R}^n} \quad & m_k(\mathbf{p}) = f_k + \nabla f_k^T \mathbf{p} + \frac{1}{2} \mathbf{p}^T \nabla^2 f_k \mathbf{p}, \\ \text{s. t.} \quad & \|\mathbf{p}\|_2 \leq \Delta_k, \end{aligned} \quad (\text{trust region})$$

without ruining **global and fast local** convergence of exact LS/TR Newton methods.

**Stopping criterion for iterative solver:**

$$\|\mathbf{r}_k\|_2^2 = \|\nabla^2 f_k \mathbf{p}_k + \nabla f_k\|_2^2 \leq \eta_k \|\nabla f_k\|_2^2,$$

where the sequence  $\{\eta_k\}$  with  $0 \leq \eta_k \leq 1$  for all  $k$  is called the **forcing sequence**.

# Local convergence

## Convergence theorem

Inexact Newton with unit steps:

- $\mathbf{x}_{k+1} = \mathbf{x}_k + \mathbf{p}_k$
- $\|\mathbf{r}_k\|_2^2 \leq \eta_k \|\nabla f_k\|_2^2$
- $0 < \eta_k \leq \eta < 1$

Then, if the starting point  $\mathbf{x}_0$  is sufficiently near  $\mathbf{x}^*$ ,

- the sequence  $\{\mathbf{x}_k\}$  converges to  $\mathbf{x}^*$ ,
- and

$$\|\nabla^2 f(\mathbf{x}^*)(\mathbf{x}_{k+1} - \mathbf{x}^*)\|_2 \leq \hat{\eta} \|\nabla^2 f(\mathbf{x}^*)(\mathbf{x}_k - \mathbf{x}^*)\|_2$$

for some constant  $\hat{\eta}$  with  $\eta < \hat{\eta} < 1$  (linear convergence).

# Local convergence

## Convergence rate

Inexact Newton with unit steps:

- $\mathbf{x}_{k+1} = \mathbf{x}_k + \mathbf{p}_k$
- $\|\mathbf{r}_k\|_2^2 \leq \eta_k \|\nabla f_k\|_2^2$
- $0 < \eta_k \leq \eta < 1$

Then,

- if  $\eta_k \rightarrow 0$ , the sequence  $\{\mathbf{x}_k\}$  converges to  $\mathbf{x}^*$  **superlinearly**;
- if  $\nabla^2 f(\mathbf{x})$  is Lipschitz continuous for  $\mathbf{x}$  near  $\mathbf{x}^*$  and  $\eta_k = O(\|\nabla f_k\|_2)$ , then the convergence is **quadratic**.

## Example:

- $\eta_k = \min(0.5, \sqrt{\|\nabla f_k\|_2})$  would yield superlinear convergence;
- $\eta_k = \min(0.5, \|\nabla f_k\|_2)$  would yield quadratic convergence.

# Line search Newton-CG method

## Algorithm

Given  $\mathbf{x}_0$ ;

**loop**

Define the forcing sequence  $\eta_k = \min(0.5, \sqrt{\|\nabla f_k\|_2})$

Use CG to solve  $\nabla^2 f_k \mathbf{p}_k = -\nabla f_k$  approximately with accuracy  $\eta_k$

Compute  $\alpha_k$  by line search method;

Update  $\mathbf{x}_{k+1} = \mathbf{x}_k + \alpha_k \mathbf{p}_k$ ;

**end loop**

**Remark:**  $\nabla^2 f_k$  is not necessarily positive definite, but the CG method is designed to solve positive definite systems. So we need to modify CG loop.

# Line search Newton-CG method

```
Given  $\mathbf{x}_0$ ;  
for  $k = 0, 1, \dots$  do  
  Set  $\mathbf{z}_0 = 0$ ,  $\mathbf{r}_0 = -\nabla f_k$ ,  $\mathbf{d}_0 = \mathbf{r}_0$  and  $\eta_k = \min(0.5, \sqrt{\|\nabla f_k\|_2})$ ;  
  for  $j = 0, 1, \dots$  do  
    if  $\mathbf{d}_j^T \nabla^2 f_k \mathbf{d}_j \leq 0$  then  
      if  $j = 0$  then  
        Stop CG and return steepest descent direction:  $\mathbf{p}_k = \mathbf{d}_0$ ;  
      else  
        Stop CG and return  $\mathbf{p}_k = \mathbf{z}_j$ ;  
      end if  
    end if  
     $\vdots$  rest of CG loop  
  end for  
  Compute  $\alpha_k$  by line search method;  
  Update  $\mathbf{x}_{k+1} = \mathbf{x}_k + \alpha_k \mathbf{p}_k$ ;  
end for
```

# Line search Newton-CG method

- Inner CG loop always produces a direction of descent for  $f$ .
- When the Hessian  $\nabla^2 f_k$  is nearly singular, the line search Newton-CG direction can be long and of poor quality.
- It does not require explicit knowledge of the Hessian, and it requires only the Hessian-vector products. Finite differencing and automatic differentiation techniques can be used.
- Preconditioning can be introduced to speed up CG.

# Limited-memory quasi-Newton methods

**Idea:** They save only a few vectors of length  $n$  that represent the approximation of the Hessian implicitly.

- Useful for solving large problems with costly or nonsparse Hessian.
- Linear convergence but fast rate.

## Limited-memory BFGS (L-BFGS):

- It uses curvature information from only the most recent  $m$  iterations to construct the Hessian approximation.
- Modest values of  $m$  ( $\sim 3 - 20$ ) work fine in practice, but the best  $m$  depends on the problem.
- Slow convergence in ill-conditioned problems.

# Limited-memory quasi-Newton methods

**Idea:** They save only a few vectors of length  $n$  that represent the approximation of the Hessian implicitly.

- Useful for solving large problems with costly or nonsparse Hessian.
- Linear convergence but fast rate.

## Limited-memory BFGS (L-BFGS):

- It uses curvature information from only the most recent  $m$  iterations to construct the Hessian approximation.
- Modest values of  $m$  ( $\sim 3 - 20$ ) work fine in practice, but the best  $m$  depends on the problem.
- Slow convergence in ill-conditioned problems.



# L-BFGS update

**Review:** BFGS inverse Hessian update:

$$H_{k+1} = V_k^T H_k V_k + \rho_k \mathbf{s}_k \mathbf{s}_k^T$$

where  $V_k = I - \rho_k \mathbf{s}_k \mathbf{y}_k^T$ ,  $\rho_k = 1/(\mathbf{y}_k^T \mathbf{s}_k)$ ,  $\mathbf{s}_k = \mathbf{x}_{k+1} - \mathbf{x}_k$  and  $\mathbf{y}_k = \nabla f_{k+1} - \nabla f_k$ .

- Since  $H_k$  is generally dense, the cost of storing and manipulating it is prohibitive when  $n$  is large.
- We store a modified version of  $H_k$  implicitly, by storing  $m \ll n$  of the vector pairs  $\{\mathbf{s}_k, \mathbf{y}_k\}$ .
- The product  $H_k \nabla f_k$  can be obtained by performing a sequence of inner products and vector summations.
- After the new iterate is computed, we replace the oldest pair with the new pair.

# L-BFGS update

Update algorithm: Compute  $H_k \nabla f_k$

```
Given  $H_k^0$ ; Set  $\mathbf{q} = \nabla f_k$ ;  
for  $i = k - 1, k - 2, \dots, k - m$  do  
     $\alpha_i = \rho_i \mathbf{s}_i^T \mathbf{q}$ ;  
     $\mathbf{q} = \mathbf{q} - \alpha_i \mathbf{y}_i$ ;  
end for  
 $\mathbf{r} = H_k^0 \mathbf{q}$ ;  
for  $i = k - m, k - m + 1, \dots, k - 1$  do  
     $\beta = \rho_i \mathbf{y}_i^T \mathbf{r}$ ;  
     $\mathbf{r} = \mathbf{r} + \mathbf{s}_i (\alpha_i - \beta)$ ;  
end for  
Output  $\mathbf{r}$ .
```

- It recursively expands the update with  $m$  pairs  $\{\mathbf{s}_k, \mathbf{y}_k\}$ .
- $H_k^0$  is allowed to vary from iteration to iteration.
- It requires  $4mn$  multiplications and calculation of  $H_k^0 \mathbf{q}$ .

# L-BFGS method

Given  $\mathbf{x}_0$  and  $m$ ;

**loop**

Choose  $H_k^0$ ;

Compute  $\mathbf{p}_k = -H_k \nabla f_k$  by update algorithm;

Compute  $\alpha_k$  by line search method;

Update  $\mathbf{x}_{k+1} = \mathbf{x}_k + \alpha_k \mathbf{p}_k$ ;

**if**  $k > m$  **then**

Discard  $\{\mathbf{s}_{k-m}, \mathbf{y}_{k-m}\}$  from storage;

**end if**

Store  $\mathbf{s}_k = \mathbf{x}_{k+1} - \mathbf{x}_k$  and  $\mathbf{y}_k = \nabla f_{k+1} - \nabla f_k$ ;

**end loop**

- A good choice for  $H_k^0$  in practice:  $H_k^0 = \gamma_k I$  with  $\gamma_k = (\mathbf{s}_{k-1}^T \mathbf{y}_{k-1}) / (\mathbf{y}_{k-1}^T \mathbf{y}_{k-1})$ .
- The line search based on the (strong) Wolfe conditions makes BFGS stable.
- The first  $m - 1$  iterates are the same as in BFGS.

## Relationship with CG methods

- Limited-memory methods historically evolved as improvements of nonlinear CG methods.
- The Hestenes-Stiefel form of nonlinear CG method:

$$\mathbf{p}_{k+1} = -\nabla f_{k+1} + \frac{\nabla f_{k+1}^T \mathbf{y}_k}{\mathbf{y}_k^T \mathbf{p}_k} \mathbf{p}_k = -\hat{H}_{k+1} \nabla f_{k+1} \quad \text{with} \quad \hat{H}_{k+1} = I - \frac{\mathbf{s}_k \mathbf{y}_k^T}{\mathbf{y}_k^T \mathbf{s}_k},$$

which resembles quasi-Newton iterates, but  $\hat{H}_{k+1}$  is neither symmetric nor positive definite.

- A symmetric positive definite modification, which also satisfies the secant equation, is

$$H_{k+1} = \left( I - \frac{\mathbf{s}_k \mathbf{y}_k^T}{\mathbf{y}_k^T \mathbf{s}_k} \right) \left( I - \frac{\mathbf{y}_k \mathbf{s}_k^T}{\mathbf{y}_k^T \mathbf{s}_k} \right) + \frac{\mathbf{s}_k \mathbf{s}_k^T}{\mathbf{y}_k^T \mathbf{s}_k},$$

which is exactly the L-BGFS method with  $m = 1$  and  $H_k^0 = I$  (*memoryless BFGS*).